

Commonplace Book - Unique Divine

Author: Unique Divine

Institute: Columbia University

Date: May, 2020 - present

Bio: <http://uniquedivine.xyz/>



Victory won't come to us unless we go to it.

Contents

1	Introduction	1
1.1	A Commonplace (Book)	1
I	Practical Reference	2
2	Metalearning & Doing	3
2.1	Ultralearning	3
2.2	Project Proposals	4
2.3	Getting Things Done (GTD)	7
3	How to	10
3.1	How to Read	10
3.2	Software Projects	11
3.3	Study 70+ hours a week efficiently	12
3.4	Notetaking	14
3.5	Learning Deeply: Knowing versus Understanding	15
II	Artificial Intelligence & Deep Learning	16
4	Deep Reinforcement Learning	17
4.1	Getting Started	17
4.2	Review Paper:	19
4.3	Mine	20
4.4	DQN	20
4.5	Collection of Refs	21
5	Attention Mechanism	22
5.1	Image + Attention	22
6	Generative Adversarial Networks	24
7	ML Finance Project	27
8	Bioinformatics	28
8.1	Deep Learning for Genomic Prediction	28
8.2	Computational Genomics (course, Rob Edwards)	32

III Math & Code	34
9 Java	35
9.1 Gradle	35
10 Git	37
10.1 Fundamental Concepts	37
10.2 Git\$2	38
10.3 Permanently removing files from commit history	38
10.4 Branching	38
11 C++	40
12 Web development	41
12.1 Hugo Web Design	41
12.2 HTTP, TCP, IP	44
13 Algorithms	47
13.1 Algorithms	47
13.2 Design Patterns	47
14 Miscellaneous	48
14.1 HTTP, REST, & Spring Boot	48
IV Journal	49
15 Quotes	50
16 Journal	52
16.1 Random	52
16.2 Algo Trading	52
16.3 Daily Schedule (when free)	53
17 2020	55

Chapter 1 Introduction

1.1 A Commonplace (Book)

A commonplace book is a system for writing down and sorting all manner of tidbits: quotes, anecdotes, observations, and information gleaned from books, conversations, movies, song lyrics, social posts, podcasts, life experiences, or anything else that you might want to return to later.

It's called a commonplace it's an aggregate, an all-encompassing system for storing all of the the most important points you learn in one common place—a central resource that makes it easy to find, re-read, and utilize each piece of wisdom you have obtained.

[Wikipedia
Common-
place Book](#)

¹

Why? | The Purpose of this Commonplace Book

1.1.1 We are at war.

Constantly seeking out new experiences and opportunities to grow. We employ active recall methods, watch videos, lectures, and seminars. We listen to mentors and podcasts, captivated by the new insights and perspectives that enter our psychology.

As voracious learners, we are in a forever war

1.1.2 Save Hours on Research

1.1.3 Find Unexpected Connections

1.1.4 Focus future Learning

¹[“What Is a Commonplace Book?” \(Article\)](#)

Part I

Practical Reference

Chapter 2 Metalearning & Doing

2.1 Ultralearning

Ultralearning: It's about learning quickly and effectively. - It's a strategy - self-directed - intense - rapid speed with great efficiency

"Despite their idiosyncrasies, the ultralearners had a lot of shared traits. They usually worked alone, often toiling for months and years without much more than a blog entry to announce their efforts. Their interests tended toward obsession. They were aggressive about optimizing their strategies, fiercely debating the merits of esoteric concepts such as interleaving practice, leech thresholds, or keyword mnemonics. Above all, they cared about learning. Their motivation to learn pushed them to tackle intense projects, even if it often came at the sacrifice of credentials or conformity."

How to become an ultralearner

1. **Metalearning:** First Draw a Map
2. **Focus:** Sharpen Your Knife
3. **Directness:** Go Straight Ahead
4. **Drill**
5. **Retrieval:** Test to Learn
6. **Feedback:** Don't Dodge the Punches
7. **Retention:** Don't Fill a Leaky Bucket
8. **Intuition: Dig Deep Before Building Up** Develop your intuition through play and exploration of concepts and skills.
9. **Experimentation**

2.1.1 Core Message (YouTube)

A lot of what's taught in school is theory useful for PhD students.

An ultralearning project is about building skills, not just knowledge.

To Start Any Ultralearning Project

1. Make and follow a Metalearning map
If I want to do X ,
 - what concepts do I need to understand?
 - what facts do I need to memorize?

- what procedures do I need to practice?

Think hard about what you want to be able to do by the end of your ultralearning project. Then identify what skills will be critical to your success.

2. Design and use practice drills

- Based on the info from metalearning, design drills.
- Practice drills should have quick feedback loops that target key areas, just like a bodybuilder would use dips to target the triceps.

3. Overlearn

Go beyond the requirements of what is necessary to learn to further reinforce what IS necessary.

To embrace overlearning, ask:

- what's my target performance, and
- what's the next level?

Commit to an even harder performance than the one you're training for.

To recap: 1. Metalearn 2. Drill 3. Overlearn

2.1.2 How to Start Your Own Ultralearning Project (Part One) [blog]

Designing your own ultralearning project has three parts:

1. Figuring out what you want to learn deeply, intensely and quickly.
2. Choosing which format you want for your project.
3. Preparing to start learning.

Step 1: What Do You Want to Ultralearn?

Step 2: Choose the Project Format

Step 3: Preparing to Learn

2.2 Project Proposals

2.2.1 PyTorch | Deep Learning

(August 25)

Structured Resources:

- [Deep Learning Book](#)
- [sentdex's PyTorch | Tutorial](#)
- [freeCodeCamp | Tutorial](#)
 - (0:00:00) Introduction
 - (0:03:25) PyTorch Basics & Linear Regression
 - (1:32:15) Image Classification with Logistic Regression

- (3:06:59) Training Deep Neural Networks on a GPU with PyTorch
- (4:44:51) Image Classification using Convolutional Neural Networks
- (6:35:11) Residual Networks, Data Augmentation and Regularization
- (8:12:08) Training Generative Adversarial Networks (GANs)
- [Intro to PyTorch | Tutorial](#)
- [MIT Deep Learning | Course](#)
- [MIT Deep Learning and AI | Lectures](#)

Keywords: Backpropagation, Hyperparameter, Feed-forward neural nets, Convolutional neural nets, Regularization

Meta Project:

There's a list of tasks below. You can mix up the steps and bounce between tutorials or start working on the project before you've finished everything else, that's completely fine. The same goes for the reading. As long as everything gets done, consider this mini project a success.

- Read through some of the deep learning book.
- Learn from sentdex's tutorial
- Complete a small project with PyTorch. Write it as a tutorial and share it with a few people.
- [freeCodeCamp | Tutorial](#)
- Gain high-level understanding of the keywords, enough that you could explain these concepts off the top of your head. Understand how they work and their significance.
- [Intro to PyTorch | Tutorial](#)

Libraries | PyTorch ignite, fastai, skorch

2.2.2 Blogging

2.2.3 Python DS Handbook

It's Tuesday, August 18th and there are about two weeks until the semester starts.

I could definitely do with improving my Python data analysis fundamentals before the semester starts. Both the Python DS Handbook and Wes Mickinney's book are awesome resources that I could use to up my skills. So here, I'm proposing a 6 day project to quickly and efficiently milk these resources.

Motivation:

I've been trailing off on too many tangents. My recent project involving LSTMs sent me on an RNN binge that turned into a neural network binge that turned into a deep learning and AI

binge. This isn't necessarily a bad thing. Some of that playful curiosity and interest was reignited in delving in this exploration. Heck, I even started reading academic papers for fun again.

Perhaps I'd lost sight of my true mission because I got too engrossed in the daily grind and tasks in front of me. Stopping to consider my interests didn't even cross my mind because lately I keep having to put out new fires. So, I took a step back. I wrote, I thought, and I've read a lot in the past few days. I burned through dozens of books and articles, particularly on finance, investing, deep learning, metalearning, and productivity. And, I'm ready to go back and apply some more of what I've learned.

Read,

Proposal:

You (I) have 6 full days to accomplish the follow, starting now:

- Read (and highlight), note-take from Python DS Handbook.
 - This is a pace of roughly 40 pages per day. Try to get through 100 a day.
- 1. **Skim.** Search for interesting content you want to learn and relevant sections. Mark or record these sections for reading.
- 2. **Read and Highlight** the sections you want to learn in more detail.
- Create content from the notes in the DS Handbook. This includes but is not limited to
 - creating anki cards
 - adding to your cookbook
- Depending on how early you finish, you can decide whether you'd gain more from going through Wes's book or some other resource.
- **Also**, study 1-2 hours of Japanese to catch up on reps each day and exercise 30 minutes minimum.

Reflection

At a high level, tell me what happened: Surprise surprise. I didn't finish the reading goal, however this does not mean that the project was a failure. I believe I discovered pretty quickly in working on this reading project that I wasn't making the most time-efficient gains from taking a bottom-up approach to NumPy.

What to change for future projects: I saw specific examples of NumPy concepts that would've helped me so much in my recent ML for Finance course. And the crazy thing was, I didn't need to know a lot of this info cold. I would've benefit just as much from having the familiarity of recognizing what utilities were even possible. Even just reading through the sections on broadcasting, masking, and universal functions would've helped me greatly.

If I could go back in time 3 months, I would read and highlight through this book cover to cover. I would take notes on some of the concepts and utilities available and potentially add

them to a cookbook instead of making Anki card active recall questions of everything.

The time saved using the text as a reference instead of a tutorial would've helped me spend more time working on end-to-end projects, where my learning environment was stimulating, motivating, and had clear transferable value to potential employers and others around me.

Key Takeaways:

1. Read first.
2. Archive.

2.3 Getting Things Done (GTD)

“If your mind is empty, it is always ready for anything; it is open to everything”-
Shunryu Suzuki

Basic workflow

1. Capture
2. Clarify
3. Organize
4. Reflect
5. Engage

2.3.1 GTD | Overview

Projects A project is any outcome for which you need more than one action step to achieve it. You often do not *do* projects, you do actions which may be related to the completion of a project.

Natural Planning of Projects is a way to think of projects to create maximum value with minimal effort and time expenditure.

1. Why | Define the purpose, principles, and values of the project
2. What | Visualize the desired outcome
3. How | Brainstorm and organize ideas on how to accomplish your goal(s)
4. When / What now? | Identify your next immediate actions.

Capture You must capture all the things you consider incomplete in your universe; personal or professional, big or small, urgent or not.

Clarify Clarifying is the act of defining and clarifying, one by one, all the stuff in your inbox, until it is empty.

- (1) Discard a task if it is worthless, (2) incubate it in the Someday/Maybe list, or (3) keep it as Reference Material, if it is potentially useful information.
- If an action takes less than two minutes, do it. If now is not the right time (be careful not to procrastinate), defer it to a specific time in the future.

paragraph*Organize

- To organize things that are actionable, you need a Projects List

- The calendar is sacred territory. If you put something there, it is because it has to be done on the date indicated.

Reflect The system cannot be static. If you want to be able to correctly choose your actions, you have to keep the system current. You need to review the whole picture of your life and your work at regular intervals and at the appropriate levels.

- Every day, you will need to review your Calendar and Next Actions list organized by contexts.
- Every week, you must review the remaining lists to keep your system clean, clear, complete and updated. The purpose of the Weekly Review is that your mind gets empty again.
- Every so often you should review the big picture; clarifying your long-term goals and visions and principles that ultimately determine how you make your decisions.

Engage

The purpose of all this management and workflow is to facilitate the decision of what you should be doing at any time.

Your priorities are set by a hierarchy of levels or perspective, from top down (1) your life purpose, (2) your vision of yourself in the future, (3) your medium-term goals, (4) your areas of responsibility, (5) your current projects and (6) the actions you need to do every day.

2.3.2 GTD | Further Reading

[The Science Behind GTD](#)

[How to find your life purpose](#)

[10 reasons why GTD will fail](#)

[The Weekly Review, in detail](#)

2.3.3 Review, in detail

Review is a key component of GTD, to the point that **if you do not do it effectively, you are probably not following the methodology the right way**. It is a time we spend to ensure that we are always working on what is actually important, a way to be proactive and take control.

Get clear

The first phase of review is to do clean-up, collecting everything and processing it. The result should be that everything is in the GTD system and the Inbox is empty. Clarify and assign dates to tasks.

Get current

The second stage is to keep your system up to date. This is **absolutely necessary for you to fully trust your system**. Review the:

- **Inbox, past, and upcoming events** | Is there anything pending? Confirm that everything is done or update deadlines if necessary.
- **Projects** | Evaluate the status of each project. Review the support material and add new actions if necessary. Make sure there is a next action for each active project.

Get creative

You have everything under control. Now it is time to dig up old longings, to imagine, to dream, and to invent.

- **Someday/Maybe list** | Is it time to activate a project that was put off or abandoned? Are there things that no longer make sense?
- **Imagine** | What are your big picture ambitions? How can you improve? What new ideas should you add to the system?

Chapter 3 How to

3.1 How to Read

3.1.1 How to read a scientific paper

2017 scientific paper by Hubbard and Dunbar, about reading scientific papers.

Hubbard, K. E., & Dunbar, S. D. (2017). Perceptions of scientific research literature and strategies for reading papers depend on academic career stage. PloS one, 12(12), e0189753.

image

1. Read the intro and background first, then the abstract (fast pass).
2. Read discussion (fast pass)
3. Attempt to interpret figures and define variables (highlight)
4. Re-read the other sections(highlight, slow pass)
5. Summarize the intro and background
6. Synthesize notes and ask questions. Record this somewhere (like a commonplace book)
7. Read methods, experiments, results, conclusion (fast pass). If necessary, go through a slow pass, highlighting and archiving the most salient points
8. **Don't forget** to take note of key references.
9. Attempt to solve all of your mysteries.

3.1.2 Three-Pass Approach - Scientific paper reading

A simple “three-pass” approach to prevent drowning in the details before drowning in the details before getting into the details. You’ll be able to adjust the depth of paper evaluation depending on your needs and available time.

The first pass gives you a general idea about the paper. The second pass lets you grasp the paper’s content, but not its details. The third pass helps you understand the paper in depth.

First pass: This pass is a quick scan to get a view from 10,000 feet and decide whether to do any more passes. It should take \approx 5 to 10 minutes. Steps:

1. Carefully read the title, abstract, and introduction.
2. Read the section and sub-section headings, but ignore everything else.
3. Glance at the mathematical content (if any) to determine the underlying theoretical foundations
4. Read the conclusions
5. Glance over the references, taking notice of which ones you’ve already read.

Second pass: Should take up to an hour. Read the paper with greater care, ignoring graphs and minute details. Note any terms you didn't understand, or questions you want to ask.

1. Understand the figures, diagrams, and other illustrations.
2. Read in full, but skip proofs and overcomplicated equations.
3. Mark relevant unread references for future reading and to learn more about the background of the paper.

You may still not understand the paper after the second pass because the subject matter is too unfamiliar. You can either (a) postpone reading the paper further until further notice, (b) return to the paper after researching the unknown terminology and questions raised, or (c) move on to the third pass.

Third pass: Virtually re-implement the paper. Identify and challenge every assumption in every statement. I called this “solving all the mysteries” in the previous section. Jot down ideas for future work and be able to explain in clear terms the entire paper from memory.

- Note source: [Keshav. 2016. How to Read a Paper.](#)
- Keywords: note, paper, reading, how to, scientific

3.2 Software Projects

3.2.1 Portfolio Overview

Key factors for a Good Portfolio

- Completeness
- Functionality
- Readability
- Documentation/information

3.2.2 Solve all of your mysteries

The time it takes programmers to finish projects varies wildly. This section will go over why some people are spending 50+ hours on projects that should be done in 5.

- Faster programmers can explain their code and strategy even if it doesn't work.
- Faster programmers dig deep and really try to understand the functions, the classes, the API's that are at their disposal so that they can apply them to solve problems.
- Slower programmers are trying to look for a shortcut. They want to use stack overflow or some other site. This difference may not be noticeable on smaller beginner projects, but over time this difference becomes huge and will ultimately be crucial to one's long-term growth and flexibility.
- Never take code on faith
 - If code is in your program, you should take time to learn about what it does

- Find some documentation. Read about the functions and classes used and DO NOT move forward until you understand each piece of code.
- Read documentation
- Solve ALL of your mysteries
 - You wrote a bunch of code, thought you knew what it did, and it's behaving mysteriously. A lot of times, people will try a couple of things. They'll sort of randomly try stuff and even if the problem goes away and they don't really know why, they just move on. This is a bad idea.
 - When you don't understand a bug, you are destined to repeat it. You need to understand why that bug occurred. Get out the documentation. Get a second pair of eyes, whatever you need to do. But, understand why this behavior is occurring.

3.3 Study 70+ hours a week efficiently

[Zach Highley | Anking](#)

Doing intellectual work efficiently is tough. It takes a lot of dedication and willpower, and yet, people aren't often strategic in their approach to being efficient. To maintain focus and efficiency with studying, work, etc., you need to think about your life as if you're a bodybuilder or Olympic athlete.

3.3.1 Sleep

- Brain gene expression during REM sleep depends on prior waking experience. Riheiro et al. 1999. [paper]: Subjects played a tower of Hanoi game during a study and were retested on the task a week later. A large improvement in the task was noted on average. However, when experimenters tampered with the REM cycle of subjects, performance saw no such improvement, even if the retest occurred just a day after the original.
- Sleep Inspires Insight. Wanger et al. 2004. [paper]: This paper displays more evidence to suggest that a night's sleep can have magical properties. A cognitive test that demanded some mathematical creativity was given to people twice with a 12 hour interval between tests. When the test subjects had a full night's rest in between tests rather than two daytime tests, the normal correct response rate more than doubled.

Q: What can we learn from studies like this?

Not only is the brain locking in knowledge for long-term storage when we sleep, it is also trying to solve problems, explore, and find ways to connect pieces of information inside our head.

There's a book called "Why We Sleep" by Matthew Walker that has a lot of helpful information on this topic. In this book, he says that once you start to drop below 7 hours of sleep, you'll start to see detriments in health and performance.

The quality of sleep is also important. Some things you can do to increase sleep quality are:

- Intensely exercise
- Increase light exposure during the day
- Reduce (blue) light exposure at night. Keep your bedroom as dark as you can and on the colder side.
- Don't consume caffeine within 6 hours of bedtime. Why? Caffeine has a fairly long half-life, which means that even after you don't feel it working, it's going to stick around in your system. The mean half-life of caffeine in the plasma of healthy individuals is about 5 hours. However, caffeine's elimination half-life may range between 2 and 10 hours.
- Avoid drugs and alcohol.

3.3.2 Exercise

- Roig et al. The effects of cardiovascular exercise on human memory. 2013. [paper]: Cardiovascular exercise performed in close temporal proximity to a session of skill practice has shown to facilitate motor memory consolidation. 2 bouts of intense treadmill running (≈ 3 min) before doing vocabulary tasks increase peripheral concentrations of catecholamines (chemical signal related to certain stresses) such as dopamine and epinephrine and accelerated the rate of learning by 20%.
- One meta-analysis (combined results of multiple scientific studies) looked at 21 studies and the effect of acute and long-term cardiovascular exercise on human memory. "Acute exercise improves memory in a time-dependent fashion by priming the molecular process involved in the encoding and consolidation of newly acquired information."

3.3.3 Environment

Where are you studying? Set a consistent place for your studying that is clean, distraction-free, and yours.

1. Take advantage of classical conditioning here. You can even just change seats when you're in work vs. play mode.
2. Clean your workspace.
3. The Life-Changing Magic of Tidying Up by Marie Kondo [book]: "It is not our memories but the person we have become because of those past experiences that we should treasure. This is the lesson these keepsakes teach us when we sort them. The space in which we live should be for the person we are becoming now, not for the person we were in the past."

3.3.4 Mindset

- Don't frame things as "I have to" but rather as "I get to". I get to learn. I get to work with these people.

- Silence the noise. Only focus on the task at hand, what you have to do right now.
- How to Stop Worrying and Start Living by Dale Carnegie [book]: “Live in day-tight compartments.” Only focus on the thing you’re going to complete today. Don’t worry about what you have to do later. Remember, don’t count the costs.
- Formulate a specific study plan for the day that is specific.

Increasing developer productivity

1. Be less random. Try to add structure to your day.
2. Stop wasting time: Cut down use of social media, analytics, etc. Turn off notifications for things.
3. Have well-planned project structure: Create a template/framework for your projects and chunk work into small tasks.
4. Use project management tools rather than email. Use your external brain rather than keeping stuff in your head.
5. Code faster.
 - Use snippets, keyboard shortcuts, vim, etc.
 - Find tools that speed up tasks.
 - Practice typing faster.
 - Learn bug trackers, testing frameworks, and other time-efficient extensions.

3.4 Notetaking

3.4.1 How to Take Effective Notes When Reading - Aun Abdi (December, 2020)

Aun has been reading a book and a half per week for the past two years. He describes his process for getting the most out of what he reads in this video.

1. Have your reading accessories ready. Make sure that you are ready for reading. Aun uses two pens and a highlighter to jot down notes both in the book and in a notebook on the side.
2. Have a question in mind before you start. In general, ask lots of questions and look into them.
3. Whilst you read, highlight and write down any ideas that resonate with you. Anything that keeps you engaged and active in the reading process is a plus. Pro tip: Make sure to read the full page/section before you highlight or take any notes. This helps prevent you from highlighting everything instead of prioritizing the important bits.
4. Some people don’t like that notetaking breaks their flow. If you find this happening, take notes after full chapters. You might find that this helps with retention because it requires extra efforts for active recall. Try writing a summary at the close of a chapter.

5. Write down some principles you want to put into action after reading every book, especially if it's a productivity book or a book about a specific skill. Aun recommends drafting 5-10 principles up with each book that fits this description.
6. Review sessions. Derived from *The 7 Habits of Highly Effective People*: Sharpening the saw. Review your notes rigorously. Aun does this every single week and schedules in time to do so. This can be as simple as transferring notes from a notebook to the computer. Schedule at least an hour every single week to review your notes.

Having a structured reading routine and structured notetaking strategy is of paramount importance for effective learning. This is key to getting the most out of what you read.

Maximizing Your Understanding of Books. R.C. Waldun. 2018.

- Bill Gates takes notes in the margins of books.
- Maintenance rehearsal versus elaborative rehearsal → You want elaborative rehearsal with your reading, new pieces of information hooking into old memories. Taking notes in the margin will help you do this.

Get the Most Out of Your Books - Be an Active Reader. The Art of Improvement. 2019. [video] Active listening: Requires understanding, full concentration, recollection of the conversation, and minimizes confusion.

By engaging with a book, arguing with the text, asking questions, and writing a lot while reading, we can perform active reading too.

- Helps with retention
- Makes transfer from book to real life more likely

Don't become a collector. Are you more concerned with the content of the book or its appearance. Remember the reason you opened a book.

Marginalia

3.4.2 Further reading:

- Man's Search For Meaning. Frankl.
- The Gulag Archipelago. Solzhenistyn.
- Astrophysics for People in a Hurry. Neil DeGrasse Tyson.

3.5 Learning Deeply: Knowing versus Understanding

Feynman: Knowing versus Understanding [video]

Part II

Artificial Intelligence & Deep Learning

Chapter 4 Deep Reinforcement Learning

4.1 Getting Started

From Lex Fridman's intro lecture on deep reinforcement learning [\[video\]](#).

Deep reinforcement learning (RL) is perhaps one of the most exciting fields in artificial intelligence (AI). It marries the beauty and power of deep neural networks to represent and comprehend the world with the ability to act on that understanding. That's basically what the creation of intelligent beings is. Recent breakthroughs in RL captivate our imagination and inspire us.

RL at high level: An agent has to make a sequence of decisions.

Types of Learning: 4 learning domains within machine learning: Supervised, semi-supervised, unsupervised, and reinforcement. People often think that supervised learning is the only domain that requires manual labeling, however in reality, all types of machine learning are "supervised" in some way by a loss function. Every type of machine learning is supervised learning to a degree. These names for the domains tell us about the cost of human labor required to obtain that supervision.

- Supervised learning: "teach by example"; RL: "teach by experience"

RL in humans: Humans appear to learn to walk (and do many other activities) through "very few examples" of trial and error. **How** is an open question. Possible answers:

- Hardware: 230 million years of bipedal movement data
- Imitation learning: Observation of other humans walking
- Algorithms: Better than backpropagation and stochastic gradient descent

Current spot: <https://youtu.be/zR11FLZ-09M?t=757>

3 Types of RL

There are countless ways to taxonomize all of the reinforcement learning algorithms. At the highest level, there are model-based algos and model-free algos. We generally put algorithms under 3 categories: model-based, value-based, and policy-based.

Model-based:

- Learn a model of the world, then plan using the model. As an agent interacts with the environment, it constructs a model for what the dynamics of the world might be.
- Better sample efficiency. Once you have a model, you can do all kinds of reasoning that doesn't require experiencing each scenario.

- In chess and in Go, the model is given to you. The agent knows the rules of the game.

Q: What is meant by “model” is model-based learning?

In this context, a model of the environment would be a function which predicts state transitions and rewards.

(cloze) AlphaZero is the reinforcement learning algorithm that learned to play Chess and Shogi.

(cloze) AlphaZero is a famous example of a model-based algorithm.

[\[AlphaZero paper, 2017\]](#)

Q: If model-based methods are more sample efficient when they work, what is the advantage of model-free methods?

Q: (cloze) Q-functions are also called action-value functions.

Q: And what’s the difference between action-value function and state-value function?

The state-value function returns the value of achieving a certain state, whereas the action-value function returns the value for choosing an action in a state.

Q: Why are Q-functions sometimes called action-value functions?

Because the Q-function gives us the value for taking an action in some state.

Note that, as of 2018, model-free methods are more popular and have been more extensively developed and tested than model-based methods.


Value-based:

- Learn the state or state-action value. Value-based methods look to estimate the quality of states and the possible actions performed in them. This quality of the state is then used to pick a “best” action.
- Act by choosing the best action in a state. This is considered indirect learning of a “policy”.

Policy-based:

- Learn the stochastic policy function that maps state to action.
- Directly learn a policy function. In other words, take as input the representation (or representation of that world) and output an action. This action will be stochastic.
- Act by sampling policy.
- Exploration is baked in. Why? The output action is stochastic, meaning it’s a r.v.

Great resource for learning about deep reinforcement learning: Open AI Spinning Up [\[link\]](#)



4.2 Review Paper:

?

Deep learning is enabling reinforcement learning (RL) to scale to problems that were previously intractable, such as learning to play video games directly from pixels. Deep reinforcement learning algorithms are also applied to robotics, allowing control policies for robots to be learned directly from camera inputs in the real world.

Topics covered: Value-based methods, policy-based methods, central algorithms in deep reinforcement learning such as the deep Q-network, trust region policy optimisation, and asynchronous advantage actor-critic. Conclude with several current areas of research within the field.

Primary goal of the field of artificial intelligence (AI) is to produce fully autonomous agents that interact with their environments to learn optimal behaviours, improving over time through trial and error. Reinforcement learning is a principled mathematical framework for experience-driven autonomous learning. RL is a general way of approaching optimisation problems by trial and error.

The most important property of deep learning is that deep neural networks can automatically find compact low-dimensional representations (features) of high-dimensional data (e.g., images, text and audio)

the use of deep learning algorithms within RL defining the field of “deep reinforcement learning” (DRL)

For a more comprehensive survey of recent efforts in DRL, including applications of DRL to areas such as natural language processing [106, 5], we refer readers to the overview by Li [78]

The first, kickstarting the revolution in DRL, was the development of an algorithm that could learn to play a range of Atari 2600 video games at a superhuman level, directly from image pixels [84]. Providing solutions for the instability of function approximation techniques in RL, this work was the first to convincingly demonstrate that RL agents could be trained on raw,

high-dimensional observations, solely based on a reward signal. The second standout success was the development of a hybrid DRL system, AlphaGo, that defeated a human world champion in Go [128], paralleling the historic achievement of IBM's Deep Blue in chess two decades earlier [19] and IBM's Watson DeepQA system that beat the best human Jeopardy! players [31].

AlphaGo was composed of neural networks that were trained using supervised and reinforcement learning, in combination with a traditional heuristic search algorithm

In a step towards even more capable agents, DRL has been used to create agents that can meta-learn ("learn to learn") [29, 156], allowing them to generalise to complex visual environments they have never seen before

One of the driving forces behind DRL is the vision of creating systems that are capable of learning how to adapt in the real world.

4.3 Mine

Q-learning

Q: (cloze) Q-learning is a model-free RL algorithm.

(cloze) Q-learning gets its name because the agent learns the quality of actions to know how to act.

(cloze) In Q-learning, "Q" function is computed to maximize expected rewards for an action taken in a given state.

(cloze) "Model-free" means that an agent does not require a model of the environment.

(cloze) For set of states S and set of actions A , the Q-function is a map s.t. $Q : S \times A \rightarrow \mathbb{R}$.

Q: In Q-learning, Q is updated according to

$$Q_{\text{new}}(s, a) := Q(s, a) + \alpha \cdot \left(r + \gamma \max_{a_b \in A} Q(s', a_b) - Q(s, a) \right)$$

.

- $Q(s, a)$: The "quality" fn. at the current state s
- α : The learning rate. $0 \leq \alpha \leq 1$. Values close to 1 make faster changes to Q .
- r : Reward received when moving from $s \rightarrow s'$
- γ : The discount factor. Quantifies how much to "discount" the
- $\max_{a_b \in A} Q(s', a_b)$: Estimate of optimal future Q-value. This would be the highest $Q(s' | a_b)$, where a_b is the "best" action and s' is the next state.

4.4 DQN

?

4.5 Collection of Refs

Deep Attention Recurrent Q-Network. 2015. ?

<https://causalai.net/r26.pdf>

[Papers with Code - RL](#)

[Papers with Code - Representation Learning](#)

[Near-Optimal Representation Learning for Hierarchical Reinforcement Learning](#)

[Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model, 2020](#)

"... deep RL algorithm based on the maximum entropy reinforcement learning framework. In this framework, the actor aims to maximize expected reward while also maximizing entropy. That is, to succeed at the task while acting as randomly as possible" - [Soft Actor-Critic: Off-Policy Maximum ENtropy Deep RL w/ a Stochastic Actor](#)

Chapter 5 Attention Mechanism

[Attention is All You Need, 2016](#)

[Rethinking Attention with Performers](#)

5.1 Image + Attention

5.1.1 Facebook AI Research applies Transformer architecture to streamline object detection models. 2020.

- Tags:
- Affiliations: Facebook AI
- [\[paper\]](#)
- [\[article\]](#)

5.1.2 Attention Agent: Neuroevolution of Self-Interpretable Agents. Tang, Ngyuen, and Ha. 2020.

- Tags: interpretability, evolutionary algorithm, transformer, self-attention, deep RL, image input
- Affiliations: Google Brain, Google Japan

5.1.3 CURL: Contrastive Unsupervised Representations for Reinforcement Learning. Srinivas et al., 2020.

- Tags: unsupervised, representation learning, deep RL, CNN, image input
- Affiliations

5.1.4 M-CURL: Masked Contrastive Representation Learning for Reinforcement Learning. Zhu et al., 2020.

- Tags: masked training, representation learning, sample efficiency, self-supervised, CNN, transformer, deep RL, BERT, contrastive learning, image input
 - Affiliations: 1. University of Science and Technology of China. 2. Microsoft Research
- Improving sample efficiency is a key research problem in reinforcement learning (RL).

Contrastive Unsupervised representations for Reinforcement Learning (CURL).

Q: Besides the involvement masked training, what's the key difference between M-CURL and CURL?

M-CURL deals with videos (seqs of images) rather than individual images. Although consecutive frames are highly correlated, CURL handles them independently. M-CURL's

main improvement, outside of getting improved performance on several benchmarks, is that it takes into consideration the correlation between sequential frames.

This is where the transformer comes in. The Transformer, together with a CNN encoder, leverages the correlation of consecutive input frames to reproduce missing features in masked frames.

Q: Why use Transformers, specifically?

The input in this paper was a sequence of images rather than a single image. Transformers (Waswani et al., 2017) are the current state-of-the-art module for modeling sequences and capturing their interdependencies.

Q: The authors call the Transformer module an "auxilliary Transformer". What makes it auxiliary?

Q: CNN encoder of what? What's being encoded? And what is meant by "encode" here?

Q: Why discard the Transformer during action selection?

Q: Policy network? What does it do? What is it made up of? What are its inputs?

Q: Contrastive learning?

Chapter 6 Generative Adversarial Networks

6.0.1 GANs ?

Abstract & Introduction

Abstract: Goodfellow et al. propose a new framework for estimating generative models via an adversarial process, in which two models are simultaneously trained:

- G : a generative model that captures the data distribution, and
- D : a discriminative model that estimates the probability that a sample came from the training data rather than G

The training procedure for G is to maximize the probability of D making a mistake. This framework corresponds to a minimax two-player game. In the space of arbitrary functions G and D , a unique solution exists, with G recovering the training data distribution and D equal to $\frac{1}{2}$ everywhere. In the case where G and D are defined by multilayer perceptrons, the entire system can be trained with backpropagation. There is no need for any Markov chains or unrolled approximate inference networks during either training or generation of samples.

Introduction: In the proposed adversarial nets framework, the generative model is pitted against an adversary: a discriminative model that learns to determine whether a sample is from the model distribution or the data distribution.

- **Counterfeiters Analogy:** The generative model can be thought of as analogous to a team of counterfeiters, trying to produce fake currency and use it without detection, while the discriminative model is analogous to the police, trying to detect the counterfeit currency. Competition in this game drives both teams to improve their methods until the counterfeits are indistinguishable from the genuine articles.

In this article, we explore the special case when the generative model generates samples by passing random noise through a multilayer perceptron, and the discriminative model is also a multilayer perceptron. We refer to this special case as adversarial nets. In this case, we can **train both models using only the highly successful backpropagation and dropout algorithms** [17] and sample from the generative model using only forward propagation.

Adversarial Net Algorithm

The adversarial modeling framework is most straightforward to apply when the models are both multilayer perceptrons. To learn the generator's distribution p_g over data \mathbf{x} , $p_g(\mathbf{x})$, we define a prior on input noise variables, $p_z(\mathbf{z})$, then represent a mapping to data space as $G(\mathbf{z}|\theta_g)$, where G is a differentiable function represented by a multilayer perceptron with parameters θ_g . We also define a second multilayer perceptron $D(\mathbf{x}|\theta_d)$ that outputs a single scalar. $D(\mathbf{x})$ represents the

probability that \mathbf{x} came from the data rather than p_g . We train D to maximize the probability of assigning the correct label to both training examples and samples from G . We simultaneously train G to minimize $\log(1 - D(G(\mathbf{z})))$. In other words, D and G play the following two-player minimax game with value function, $V(G, D)$:

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log (1 - D(G(\mathbf{z})))].$$

Goodfellow et al. present a theoretical analysis of adversarial nets, essentially showing that the training criterion allows one to recover the data generating distribution as G and D are given enough capacity, i.e., in the non-parametric limit.

Figure 1 explains this approach: GANs are trained by simultaneously updating the discriminative distribution, D , so that it discriminates between samples from the data generating distribution, p_x , and those of the generative distribution p_g (G).

After several steps of training, if G and D have enough capacity, they will reach a point at which both cannot improve because $p_g = p_{\text{data}}$. The discriminator is unable to differentiate b/w the two distributions, i.e. $D(\mathbf{x}) = \frac{1}{2}$.

Algorithm 1: Minibatch SGD training of GAN. The number of steps to apply the discriminator, k , is a hyperparameter. $k = 1$ is the least expensive option.

for number of training iterations **do**

for k steps **do**

 # Apply discriminator

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior, $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution, $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log \left(D(\mathbf{x}^{(i)}) \right) + \log \left(1 - D(G(\mathbf{z}^{(i)})) \right) \right].$$

end

 # Update generator

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior, $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left(1 - D(G(\mathbf{z}^{(i)})) \right).$$

end

The gradient-based updates can use any standard gradient-based learning rule.

Note:

- $k = 1$ was used in the experiments.
- Momentum was the gradient-based learning rule used in the experiments.

Related Work

The adversarial nets framework does not require a Markov chain for sampling. Because adversarial nets do not require feedback loops during generation, they are better able to leverage piece-wise linear units [19, 9, 10], which improve the performance of backpropagation but have problems with unbounded activation when used in a feedback loop.

?

Chapter 7 ML Finance Project

example w/ multivariate time series in PyTorch

Q: (cloze) Neural networks can be constructed using the `torch.nn` package.

Q: Import the package for constructing neural networks in PyTorch.

```
import torch.nn as nn
```

Q: (cloze) Seaborn comes with built-in datasets.

Q: Load seaborn's flights dataset.

```
flight_data = sns.load_dataset("flights")
```

Q: Why must time series data be scaled for sequence predictions?

When a network is fit on unscaled data, it is possible for large inputs to slow down the learning and convergence of your network and in some cases prevent the network from effectively learning your problem.

Q: sklearn import for scaling data?

```
from sklearn.preprocessing import MinMaxScaler
```

We know the field is fast moving. If the reader looking for more recent free reading resources, there are some good introductory/tutorial/survey papers on Arxiv; I happen to be compiling a list of them.

One of said review papers ?

(?, hello)

Chapter 8 Bioinformatics

8.1 Deep Learning for Genomic Prediction

original title: DL for Genomic Risk Scores

“ A central aim of computational genomics is to identify variants (SNPs) in the genome which increase risks for diseases. Current analyses apply linear regression to identify SNPs with large associations, which are collected into a function called a Polygenic Risk Score (PRS) to predict disease for newly genotyped individuals. This project is broadly interested in whether we can improve performance of genomic risk scores using modern machine learning techniques.

A recent study assessed the disease prediction performance of neural networks in comparison to conventional PRSs, but did not find evidence of improvement. This project will explore whether neural networks can improve performance by incorporating gene expression data to the training process. Gene expression is often integrated with SNP data in Transcriptome-Wide Association Studies (TWAS), which bear some resemblance to neural network architectures with SNPs as input nodes, genes as intermediate nodes, and disease status as the output node. Modeling this process as a neural network however will require defining a more unconventional architecture in which a small subset of hidden nodes is anchored to observed values.

This project is designed for students with experience in machine learning topics and preferably with deep learning tools such as tensorflow or pytorch. Students should also be interested in applying machine learning and statistics to genomics applications.” - Jie Yuan

Terms to know: Computational genomics, variants, single-nucleotide polymorphism (SNP), genome, Polygenic Risk Score (PRS), Transcriptome-Wide Association Studies (TWAS), gene(s), genomics, genotype, intermediate node (NN), hidden node (NN)

Q: Explain at a high level how PRSs tells us disease risk.

Q: “In what task were neural networks outperformed by conventional PRSs?”

Q: What is meant by “improve performance of genomic risk scores?”

genotype : An individual’s collection of genes. Also can refer to the two alleles inherited for a particular gene.

node (NN) : An artificial neural network is an interconnected group of nodes, inspired by a simplification of neurons in a brain. Here, each circular node represents an artificial neuron and an arrow represents a connection from the output of one artificial neuron to the input

of another¹.

hidden node (NN) : A node in a hidden layer.

hidden layer (NN) :

8.1.1 Papers

Polygenic Risk Scores (paper) ?

Abstract (mining)

recurrence risks : In genetics, the likelihood that a hereditary trait or disorder present in one family member will occur again in other family members².

“Evidence for genetic contribution to complex diseases is described by recurrence risks to relatives of diseased individuals.”

This is distinguished from recurrence risk for cancer, which is the chance that a cancer that has been treated will recur.

gene : a sequence of DNA that codes for a specific peptide or RNA molecule; the physical and functional unit of heredity.

locus : the position of a gene on a chromosomes

somatic cell : any cell of the body except sperm and egg cells. A non-germline cell. any biological cell forming the body of an organism (except gametes).

sôma (Ancient Greek): body

genome : An organism’s complete set of DNA, including all of its genes. Each genome contains all of the information needed to build and maintain that organism. In humans, a copy of the entire genome—more than 3 billion DNA base pairs—is contained in all cells that have a nucleus ³.

“genome-wide association”

allosome : (1) A sex chromosome such as the X and Y human sex chromosomes. (2) An atypical chromosome ⁴.

allo- (Greek): other, different

autosome : Any chromosome that is not a sex chromosome. The numbered chromosomes.

auto (Greek): self, one’s own, by oneself, of oneself

-some, soma (Greek): body

allele : (genetics) One of a number of alternative forms of the same gene occupying a given position, or locus, on a chromosome.

Borrowed from German Allel, shortened from English allelomorph. Ultimately from the Ancient Greek prefix allēl- from állos (“other”).

¹https://en.wikipedia.org/wiki/Artificial_neural_network

²<https://www.cancer.gov/publications/dictionaries/genetics-dictionary/def/recurrence-risk>

³<https://ghr.nlm.nih.gov/primer/hgp/genome>

⁴<https://www.merriam-webster.com/medical/allosome>

“their effects and allele frequencies”

allelomorph: another term for allele.

risk loci :

“genome-wide association studies allow a description of the genetics of the same diseases in terms of risk loci...”

haploid : the quality of a cell or organism having a single set of chromosomes.

diploid : the quality of having two sets of chromosomes.

“Sexually reproducing organisms are diploid” (having two sets of chromosomes, one from each parent)

eukaryotes : Organisms whose cells have a nucleus enclosed within a nuclear envelope.

gamete : A mature sexual reproductive cell, as a sperm or egg, that unites with another cell to form a new organism. A haploid cell that fuses with another haploid cell during fertilization in organisms that sexually reproduce. A mature haploid male or female germ cell which is able to unite with another of the opposite sex in sexual reproduction to form a zygote.

gamete (Ancient Greek): to marry

zygote : A eukaryotic cell formed by a fertilization event between two gametes.

zygōtos (Greek): joined. yoked.

monozygotic : Monozygotic (MZ) or identical twins occur when a single egg is fertilized to form one zygote (hence, "monozygotic") which then divides into two separate embryos.

“monozygotic twins”

empirical :

“generate results more consistent with empirical estimates”

genetic variants :

Q: A human cell containing 22 autosomes and a Y chromosome is a sperm.

Neural Networks for Genomic Prediction (paper) ?

Transcriptome Wide Association

8.1.2 Meetings

Jie Yuan meeting #2 (Sep 10)

- There's a dot product and its output is passed through some activation function like sigmoid or ReLU. It's still linear in the sense that there's some sort of function that takes in a dot product (which is the definition of a linear operation). You can think of [the betas] as the weights in the neural network because the betas are the coefficients of a linear model.
- Liability is discussed in the multi-locus models paper.
- In the genomics context, this is called the **liability threshold model**. If you google that, you'll find some papers. In a broader machine learning sense, this is called the **probit regression model**. Wikipedia probably has a sufficient article on it.

- The gist of it is that it's a model to map continuous sums (the dot products basically) into binary labels: cases and controls. This can be any sort of logistic regression type thing where you have 0s and 1s. Basically, the **probit model is an alternative to the logistic regression model**.
- The liability is related to what's called the "link function". The link function in the probit model is the normal CDF function. There's a term that goes into the link function, and that term is what we're calling liability.
- The liability is basically the product of the vector of genotypes and the vector of betas. It's the linear part of the generalized linear model.
- Once you have the liability and plug it into the link function. Here, the link function is the probit regression. In a neural network, this would be the linear activation function (sigmoid, ReLU, etc.).
- You know how a sigmoid has a domain that's $(-\infty, \infty)$? It's range is $(0,1)$, so what that function does is map a real value into being a probability. It gives you the probability of being a case. If you look at the normal CDF function, you find that it looks almost the same.

Why is the liability threshold model better than or different from the logistic reg model?

- One reason people use log reg more often is that the betas from log reg are interpretable. They have a meaning in terms of odds ratios. In a log reg model, the effect sizes are the natural log of the odds ratios.
- One disadvantage of log reg compared to liability-threshold is that log reg doesn't have a concept of the underlying distribution of the liability. In log reg, you get the linear $X\beta$ term, but there's no sort of distribution around it. In the liability-thresh model, we say that the liability, itself, is standard normally distributed. By taking the normal CDF, what you're doing is mapping that $X\beta$ value, the liability value, onto this distribution and asking, "what's the probability that it's larger than some threshold?"
- larger than threshold \implies label as 1, smaller \implies 0.
- Because liability-threshold has that normal distribution, it gives you more concepts to play with such as variance explained, which the logistic reg model doesn't have.
- [liability threshold model in simple terms](#)
from wikipedia: [Liability threshold model](#)

schizophrenia (SZ) example:

- Whether or not someone has SZ comes from a combination of factors in both their genes and environment.
- There are all sorts of variables. Imagine hypothetically that you could collect all of them and produce a score from that (that determines whether or not you have SZ). That score

would be the liability.

- The liability is a $\mathcal{N}(0, 1)$ distribution that, if higher than the threshold, predicts/indicates the patient has SZ.
- The problem with that is that we don't actually observe everything. We don't truly observe every factor that goes into whether you have SZ b/c (1) we can't measure one's environment and (2) we can't know all of the genomic risk factors from GWAS. We can only identify the largest risk factors.

This means we are identifying only a small fraction of what goes into genetic risk for SZ.

- Essentially, what we have is a small subset of the factors that actually determine your SZ risk (which we're reasonably confident indicate this relationship). In this example, some subset of the risk factors like the genomic variance which increase your risk for SZ.
- What that means is that inside this standard normal liability distribution/function, what we actually know is a small subset of the factors that go into that. If we score people on that subset of factors, what we get is a smaller distribution, a distribution that has a variance that's a small fraction of the total liability variance, which we can't observe. The variance of that small fraction that we observe, that's the **variance explained**.
-

8.2 Computational Genomics (course, Rob Edwards)

This section details my notes from [Rob Edwards's open source course](#) at San Diego State University (Copyright 2018).

What will be learned in this course?

We'll use cutting-edge tools to analyze microbial genomes.

By the time we're done with the course, you'll have the ability to:

- use Amazon Web Services to analyze genomes
- use existing bioinformatics applications
- describe how algorithms used in bioinformatics work
- download, identify, and analyze data from public repositories
- critically analyze genomes and metagenomes

8.2.1 Central Dogma of Biology

[\[lecture video\]](#)

Dogma (def):

1. a principle or set of principles laid down by an authority as incontrovertibly true."the Christian dogma of the Trinity". "the rejection of political dogma." "the classic dogma of objectivity in scientific observation". "the difficulty of resisting political dogma".

2. Characterized by assertion of unproved or unprovable principles. A doctrine that is proclaimed as true without proof. "he believed all the Marxist dogma". "dogmatic writings".

In Essence: DNA is the genetic code. DNA is converted to messenger RNA (mRNA) through a process called transcription. There are two other types of RNA: tRNA and rRNA. mRNA is converted into proteins. Proteins are comprised of amino acids. In mRNA, nitrogenous bases constitute what's called a codon. A codon encodes for one amino acid. mRNA is read sequentially, one codon at a time, to give sequences of amino acids that make up a protein. [bonus] There's also a "reverse" transcription that certain viruses can do, where mRNA is converted back into DNA. Humans don't normally do that. Bacteria don't normally do that. Viruses do. **Bottom Line:** DNA is essentially the standard code for all living organisms as far as we know.

DNA's alphabet: A C G T

RNA's alphabet: A C G U

DNA gets transcribed into RNA in a sequence-dependent fashion.

DNA has a direction. One side of a strand is 5', "the 5 prime band", and one is 3'. This direction is 5' to 3'.

DNA lines up in pairs of sequences with bases aligned in complementary pairs. These pairs are called "reverse compliments"

RNA is made with what's called a template strand.

8.2.2 What does it mean to sequence a genome?

Q: What's a genome?

Humans have 23 pairs of chromosomes. Each pair consists of a chromosome from each parent.

All of the DNA contained in one cell is called the genome. We have one copy of the genome in nearly every cell in our body. Human genomes are $\approx 99.8\%$ identical to that of every other human being. The other 0.2% of the genome is what is of high interest to healthcare professionals as understanding it can help in the prediction, prevention, diagnosis, and treatment of disease.

Q: e

- Human genome 3.1 billion bp (base pairs), i.e. 3.1 Gbp
- bacteria 100 kbp - 2Mbp
- This means that the computational overhead of studying bacterial genomes is much smaller and can be done on a standard personal computer.

Part III

Math & Code

Chapter 9 Java

Let's dissect the following code block that sums two numbers.

```
import java.util.*;
public class Solution {
    public static void main(String[] args) {
        Scanner in = new Scanner(System.in);
        int a, b;
        a = in.nextInt();
        b = in.nextInt();
        System.out.println(a + b);
    }
}
```

Q: What does `String[] args` mean in the hello world program?

`args` is an array of strings.

Arrays

Q: Declare an array of strings.

```
String[] arr;
```

Q: Initialize an array of integers containing 1, and 2.

```
int[] arr = {1, 2};
```

Q: Initialize an array of strings containing "BMW" and "Ford".

```
String[] carBrands = {"BMW", "Ford"};
```

Q: Given an array, `String[] carBrands = {"BMW", "Ford"};` print the first element.

```
System.out.println(carBrands[0]);
```

Q: `in.nextInt()` ?

content

It may be a good idea to go through all of this: <https://www.w3schools.com/java/default.asp>

9.1 Gradle

<https://youtu.be/aYu994I8Z6I?list=PLMxpKvJf0K0QUyvmkKZu7WpwTVzdePGT1>

- Gradle is an open-source build automation tool.
- Gradle is the official Android build tool; maintained by Android SDK Tools team
- Gradle build scripts are written using Groovy

WHy Gradle?

- highly customizable and extensible
- used for multiple languages
- it's fast: reuses outputs from previous executions, processing only inputs that changed; executes tasks in parallel
- higher performance than its peers (such as Maven)

Gradle project tasks

Chapter 10 Git

Git (cookbook)

10.1 Fundamental Concepts

10.1.1 Local branch vs. remote branch

- **local branch:** a branch only the local user can see. It exists only on your local machine.
 - Ex. Create local branch named "myNewBranch":

```
git branch myNewBranch
```
- **remote branch:** a branch on a remote location (in most cases 'origin'). Local branches can be pushed to 'origin' (a remote branch), where other users can track it.
 - Ex. Push local branch, "myNewBranch", to the remote, "origin" so that a new branch named "myNewBranch" is created on the remote machine ("origin"):

```
git push -u origin myNewBranch
```

- **remote tracking branch:** A local copy of a remote branch.

When 'myNewBranch' is pushed to 'origin' using the command above, a remote tracking branch named 'origin/myNewBranch' is created on your local machine.

local tracking branch: a local branch that is tracking another branch.

source(s): [SNce & Brian Webster](#). stackoverflow.com

10.1.2 HEAD, master, and origin

I highly recommend the book "Pro Git" by Scott Chacon. Take time and really read it, while exploring an actual git repo as you do.

- **HEAD:** the current commit your repo is on. Most of the time HEAD points to the latest commit in your current branch, but that doesn't have to be the case. HEAD really just means "what is my repo currently pointing at".

In the event that the commit HEAD refers to is not the tip of any branch, this is called a "detached head".
- **master:** the name of the default branch that git creates for you when first creating a repo. In most cases, "master" means "the main branch". Most shops have everyone pushing to master, and master is considered the definitive view of the repo. But it's also common for release branches to be made off of master for releasing. Your local repo has its own master branch, that almost always follows the master of a remote repo.
- **origin:** the default name that git gives to your main remote repo. Your box has its own repo, and you most likely push out to some remote repo that you and all your coworkers push to. That remote repo is almost always called origin, but it doesn't have to be.

- `HEAD` is an official notion in git. `HEAD` always has a well-defined meaning. `master` and `origin` are common names usually used in git, but they don't have to be.

source: [HEAD, master, and origin](#). Matt Greer & Jacqueline P. via [stackoverflow.com](#)

10.1.3 Large File Storage

<https://git-lfs.github.com/>

10.2 Git§2

10.2.1 SSH keys

An SSH key is an alternative to username/password authorization on GitHub. This will allow you to bypass entering your username and password for future GitHub commands.

SSH keys come in pairs, a public key that gets shared with services like GitHub, and a private key that is stored only on your computer. If the keys match, you're granted access.

The cryptography behind SSH keys ensures that no one can reverse engineer your private key from the public one.

[SSH Keys for GitHub \[article\]](#)

Generating a new SSH key: Follow [Generating a new SSH key and adding it to the ssh-agent \[article\]](#)

10.3 Permanently removing files from commit history

WHy do this? You may have committed a password, some other sensitive information, or a large file that you want to remove from github. If the change is only a few commits back, you can "rebase" changes out of the history. I actually need to do this for a much older set of files and accidentally uploaded a textbook that takes up almost a GB of space.

```
git filter-branch -force -index-filter "git rm -cached -ignore-unmatch PathToSensitiveFile" --prune-empty --tag-name-filter cat --all
```

All you need to change is the `PathToSensitiveFile` item. Once you've used this command for all of the files you'd like to get rid of, update the origin by typing `git push origin -force -all`.

10.4 Branching

Suppose your application is stable. Later, you discover a gigantic bug that was passing silently. You want to write some tests, fix the bug, and eventually have a stable, passing application once again. To do this, you'd create a branch for the fix and push the branch to the remote so that all of the developers on your team can collaborate and make the fix.

Once all of the necessary changes have been made and the application is stable, someone from the team would commit merge the commits from the other branch into master. Since the commit history from the branch will have been saved to master, the new branch could be deleted without loss of information (if you no longer wanted to work on this branch).

View branches: `git branch`

Switch branches: `git checkout [branch-name]`

10.4.1 Merging

Merge the specified branch's history into the current one. `git merge [branch]`

10.4.2 Deleting branches

Delete local branch `git branch -d [branch-name]`

Delete remote branch `git push origin -delete [branch-name]`

Chapter 11 C++

C++ source code files end with a .cpp extension.

Hello world program: Run these and find out which one works.

```
#include <iostream>

int main()
{
    std::cout << "Hello, world!";
    return 0;
}
```

Compiling and executing the C++ program:

1. Step 1 is to install the gcc compiler.

hi

2. Verify the install of g++ and gdb with `whereis g++` and `whereis gdb`

To install gdb (linux or WSL), use `sudo apt-get install build-essential gdb`

Open MP To import:

```
#include <omp.h>
```

People use OpenMP for shared memory parallelization.

Header files

C++ programs consist of more than just .cpp files. They also use **header files**, which can have a .h extension, .hpp extension, or even none at all.

Q: What is a .h file?

header file

Q: What is the purpose of a header file?

Header files allow us to put declarations in one location and then import them wherever we need them. This can save a lot of typing in multi-file programs.

Q: What's contained in a .h file?

..

References & Further Reading

- [C header files](#)
- learncpp.com/.../header-files

Chapter 12 Web development

12.1 Hugo Web Design

Start Sep 16 (Mike Dane Tutorial Series)

12.1.1 Intro to Hugo -(video)

- Hugo is a static site generator.
- Static website generators allow you to compromise between writing a bunch of static html pages and using a heavy, and potentially expensive, content management system.
- Why Hugo? It's extremely fast.
- 2 kinds of websites, dynamic and static. Dyanmic ex. Facebook. Facebook pages are dynamically generated for each user. For static websites, what you see is what you get.
- Static websites are notoriously harder to maintain b/c you lack some of the flexibility of things on a dynamic site. Usually you can't use much conditional logic, functions, or variables.
- However, static pages are extremely fast.
- Hugo is great for a blog, portfolio website, etc.
- Hugo doesn't explicity require you to write a single line of HTML code.
- Flexibility | With that said, if you want to go in and change every little detail of the layout of the site, you cna do that. You can write as much of the HTML and have as much control as you'd like.
- Hugo is 100% free and open-source.

12.1.2 Intalling Hugo on Windows - (video)

- Mine's already installed. I'll skip this for now.
- a
- a

12.1.3 Creating a new site - (video)

- skip, a bit too easy

12.1.4 Installing & Using Themes - (video)

- my themes are already installed

12.1.5 Creating & Organizing Content - (video)

- Hugo has 2 types of content: single pages and list pages
- List content lists other content on the site. You can call this a list page.
- Individual blog posts are single pages.
- Your posts should not just be in the content directory. They should be in directories inside the content directory.
- A list page is automatically created for directories inside the content folder. Hugo automatically does this. Note, this only occurs for directories at the “root” level of the content directory. For example: `content/post/` would generate a list page at `site.com/post/`, but `content/post/dir0` would not.
- If you want a list page to be generated for a dir that is not at the root level of the content dir, you have to create an **index file**, `_index.md`. For a convenient and efficient way to do this from the cmd, use `hugo new post/dir0/_index.md` (above above example), then there will be a list page for `dir0`. Content can also be added to `_index.md` and it should show up on the page.
- Additionally, for list pages that are automatically generate by hugo, you can edit the content by adding an `index.md` to those as well. Ex. `content/post/_index.md`.

12.1.6 Front Matter- (video)

- Front matter in Hugo is what is commonly called meta data.
- Front matter is data about our content files.
- The metadata automatically generated by Hugo at the top of md files when using `hugo new` is front matter
- Front matter is stored in key-value pairs
- Front matter can be written in 3 different languages: YAML, TOML, and JSON
- The default lang for front matter in Hugo is YAML
- YAML - indicated by “-”,
- TOML - indicated by “+++” and uses “=” instead of “:”,
- JSON - indicated
- You can create your own custom front matter variables.
- Front matter is super powerful in its utility.

12.1.7 Archetypes - (video)

- How does the default front matter from using `hugo new` .md get selected? Short answer: archetypes
- An archetype is basically the default front matter template for when you create a new content file.

- Archetypes are modified under `static/themes/archetypes/default.md`
- Suppose your content dir has a subdirectory, `content/dir0`. If you wanted to create an archetype for the files in `dir0`, you'd simply create `dir0.md` inside the `archetypes` dir.

12.1.8 Shortcodes - (video)

- Shortcodes are predefined chunks of HTML that you can insert into your markdown files.
- Let's say you have a md file that you want to spice up by adding in some custom HTML. For instance, maybe you'd like to embed a YouTube video. Normally this would require lots of HTML that you'd have to paste it. Shortcodes can allow you to sidestep this. Hugo comes with a YouTube video shortcode predefined.
- General shortcode syntax `< shortcode_name param0 >`
- Youtube shortcode | For a YouTube video with url, "youtube.com/watch?v=random-text", the shortcode we'd use to embed would be `< youtube random-text >` because "random-text" is the id of the youtube video and the only parameter for that shortcode.

12.1.9 Taxonomies - (video)

- Taxonomies in hugo are basically ways that you can logically group different pieces of content together in order to organize it in a more efficient way.
- Hugo provides 2 default taxonomies: tags & categories
- All taxonomy information is declared in front matter. In YAML, tags has the syntax `tags: ["tag0", "tag1", ...]`

12.1.10 Templates - (video)

- Templates here mostly refers to HTML templates. If you're not comfortable writing HTML, CSS, and coding for the web, templates might be a little bit above your head.
- A hugo theme is actually made up of hugo templates.
- Any template that you use in Hugo is going to be inside `themes/theme-name/layouts`. This is where all the templates live.
- `/layouts/default` usually contains a default style for list and single pages by use of `list.html` and `single.html`.

12.1.11 List Templates - (video)

- List templates give default HTML layout to list content files.
- .

Resources

- [A clear and concise beginner hugo tutorial](#)

- CSS Crash Course for Absolute Beginners

12.2 HTTP, TCP, IP

12.2.1 HTTP

Q: (cloze) HTTP = Hypertext Transfer Protocol

Q: Give me a brief overview of HTTP.

HTTP is a protocol that allows the fetching of resources such as HTML documents. It allows web-based apps to communicate and exchange data.

HTTP is a client-server protocol. This means requests are initiated by a recipient, usually the web browser, and a complete document is constructed from the files fetched such as text, layout description, images, videos, scripts, etc.

- The client is the one making the request.
- The server responds to this request.

Three important things about HTTP

1. HTTP is connectionless: After making the request, the client disconnects from the server. Then when the response is ready, the server re-establishes the connection again and delivers the response.
2. HTTP can deliver any sort of data.
3. HTTP is stateless: The client and server know about each other only during the current request. If the current request closes and the two computers want to connect again, they need to provide information to each other anew. In other words, statelessness means that the connection between the browser and the server is lost once the transaction ends.

To learn about requests and responses: <https://youtu.be/eesqK59rhGA?t=275>

12.2.2 TCP/IP

The internet protocol suite is the conceptual model and set of communications protocols used in the internet and similar computer networks. It is commonly known as TCP/IP because the foundational protocols in the suite are the Transmission Control Protocol (TCP) and Internet Protocol (IP).

Communication protocol: A system of rules that allow two or more entities of a communications system to transmit information via any kind of variation of a physical quantity.

Internet Protocol (IP): IP has the task of delivering packets from the source host to the destination host solely based on the IP addresses in the packet headers. For this purpose, IP defines packet structures that encapsulate the data to be delivered. It also defines addressing methods that are used to label the datagram with source and destination information.

- (network) packets: Formatted units of data carried by a packet-switched network. A packet consists of control information and user data; the latter is also known as the payload. Control information provides data for delivering the payload (e.g. source and destination network addresses, error detection codes, or sequencing information).

Q: What's an IP address?

An Internet Protocol address, or IP address, is a numerical label assigned to each device connected to a computer network that uses the Internet Protocol for communication.

Bandwidth (computing): (Computing) bandwidth is the maximum rate of data transfer across a given path. Bandwidth can be characterized as network bandwidth, data bandwidth, or digital bandwidth.

Computing bandwidth is different from the bandwidth defined in the field of signal processing, signal bandwidth. Signal bandwidth is the frequency range between lowest and highest attainable frequency while meeting a well-defined impairment level in signal power. It's measured in hertz.

In relation to internet protocol, we often talk about consumed bandwidth in bit/s, which corresponds to achieved throughput or goodput, i.e. the avg rate of successful data transfer through a communication path.

Q: Why is network bandwidth an average rate instead of a current rate?

A channel with v_β may not necessarily transmit data at v_t rate since protocols, encryption, and other factors can add overhead. For instance, internet traffic often uses TCP, which requires a three-way handshake for each transaction. TCP is efficient, but it does add significant overhead compared to simpler protocols. Additionally, data packets may be lost, further reducing the data throughput.

Q: Packet loss?

Packet loss occurs when one or more packets of data travelling across a computer network fail to reach their destination. Packet loss is either caused by errors in data transmission, typically across wireless networks, or network congestion. Packet loss is measured as a percentage of packets lost with respect to packets sent.

Q: Packet switching?

Packet switching is a method of grouping data that is transmitted over a digital network into packets. Packets are made of a header and a payload. Data in the header is used by networking hardware to direct the packet to its destination, where the payload is extracted and used by application software. Packet switching is the primary basis for data communications in computer networks worldwide.

Q:

Resources

- https://en.wikipedia.org/wiki/Internet_protocol_suite
- .

Chapter 13 Algorithms

13.1 Algorithms

"4.5 years of learning programming and working as fullstack software engineer ... had interview with one of the FAANG companies this summer in Hong Kong but failed it due to the fact that I suck in DSA (Data Structures & Algorithms)."

"I'm using to leetcode.com to learn data structures and algorithms since I got a rejection from FAANG after interviewing with them onsite."

[Role of DSA in Programming \(July, 2020\)](#)

13.2 Design Patterns

Q: Why use "design patterns"?

- Design patterns let you write better code more quickly by providing a clearer picture of how to implement the design
- Design patterns encourage code reuse and accommodate change by supplying well-tested mechanisms for delegation, composition, and other non-inheritance based reuse techniques
- Design patterns encourage more legible and maintainable code

Q: Delegation? Composition?

- delegation: a pattern where a given object provides an interface to a set of operations. However, the actual work for those operations is performed by one or more other objects.
- composition: Creating objects with other objects as members. Should be used when a "has-a" relationship appears.

Q: What are design patterns?

content

Q: Which resources will you use to start learning about design patterns?

GOF patterns (C++). Then, potentially Head First Design Patterns (Java)/

13.2.1 References & Further Reading

[Introduction to Design Patterns Course](#)

Chapter 14 Miscellaneous

14.1 HTTP, REST, & Spring Boot

Q: (cloze) REST has become the de-facto standard for building web services. Why?
Easy to build, easy to consume.

Q: What is a micro service and why do we use them?

Microservice is an architecture that allows the developers to develop and deploy services independently. Each service running has its own process and this achieves the lightweight model to support business applications.

Microservices, a.k.a. the microservice architecture, structure an application as a collection of services that are

- Highly maintainable and testable
- Loosely coupled
- Independently deployable
- Organized around business capabilities
- Owned by a small team

Q: What is Spring Boot?

An open source Java-based framework used to create microservices.

Q: Why use Spring Boot?

- manages REST endpoints
- flexible way to configure database transactions
- auto configured as opposed to manually configured
- eases dependency management

Q: How does Spring Boot easy dependency management?

Spring Boot automatically configures your application based on the dependencies you have added to the project by using `@EnableAutoConfiguration` annotation. For example, if MySQL database is on your classpath, but you have not configured any database connection, then Spring Boot auto-configures an in-memory database.

The entry point of the spring boot application is the class contains `@SpringBootApplication` annotation and the main method.

14.1.1 References & Further Reading

[Building REST services with Spring.](#)

[What are microservices?](#)

[Spring Boot - Introduction](#)

Part IV

Journal

Chapter 15 Quotes

“your mind is for having ideas, not holding them.” - David Allen

“Some very specific but seemingly mundane behaviors, when applied, produce the capacity to exist in a kind of sophisticated spontaneity, which, in my experience, is a key element to a successful life.” - David Allen

“We should hunt out the helpful pieces of teaching and the spirited and noble-minded sayings which are capable of immediate practical application—not far-fetched or archaic expressions or extravagant metaphors and figures of speech—and learn them so well that words become works.” - Seneca

“I have a friend who’s an artist and has sometimes taken a view which I don’t agree with very well. He’ll hold up a flower and say “look how beautiful it is,” and I’ll agree. Then he says “I as an artist can see how beautiful this is but you as a scientist take this all apart and it becomes a dull thing,” and I think that he’s kind of nutty. First of all, the beauty that he sees is available to other people and to me too, I believe. . .

I can appreciate the beauty of a flower. At the same time, I see much more about the flower than he sees. I could imagine the cells in there, the complicated actions inside, which also have a beauty. I mean it’s not just beauty at this dimension, at one centimeter; there’s also beauty at smaller dimensions, the inner structure, also the processes. The fact that the colors in the flower evolved in order to attract insects to pollinate it is interesting; it means that insects can see the color. It adds a question: does this aesthetic sense also exist in the lower forms? Why is it aesthetic? All kinds of interesting questions which the science knowledge only adds to the excitement, the mystery and the awe of a flower. It only adds. I don’t understand how it subtracts.”
- Richard Feynman on the beauty of a flower

“AI began with an ancient wish to forge the gods.” - Pamela McCorduck, *Machines Who Think*, 1979

“We cannot solve our problems with the same thinking we used when we created them.” - Albert Einstein

“I can calculate the motion of heavenly bodies, but not the madness of people.” - Isaac Newton

“Everyone considers changing the world, but no one considers changing himself.” - Leo Tolstoy, *Author of War and Peace*.

“Knowing yourself is the beginning of all wisdom.” - Aristotle

“Games are won by players who focus on the playing field, not by those whose eyes are glued to the scoreboard.” - Warren Buffett

“We should not only use all the brains we have but all that we can borrow.” - Woodrow Wilson

“This contest, Brian is biggest competitor because this year we are very close. I must look what he do, but still I looking what I do and if I be better than me, I will be winner. ” - Žydrūnas Savickas [\[video\]](#)

“Nothing else matters except sitting down everyday and trying. Why is this so important? Because when we sit down day after day and keep grinding, something mysterious starts to happen. A process is set into motion by which, inevitably and infallibly, heaven comes to our aid.” - Steven Pressfield, The War of Art

Chapter 16 Journal

16.1 Random

Who is François Chollet?

- works on deep learning at Google in Mountain View, CA.
- creator of the Keras deep-learning library
- contributor to the TensorFlow machine-learning framework.
- wrote a book on deep learning that is very applied rather than theoretical.

“We know the field is fast moving. If the reader looking for more recent free reading resources, there are some good introductory/tutorial/survey papers on Arxiv; I happen to be compiling a list of them”. - Auto Snipe

One of said review papers ?

16.2 Algo Trading

16.2.1 ML Finance Project

Terms to Learn: NLP, bag of words, stop words, stemming, tokenizing, term-frequency inverse document frequency, bigrams, unigrams, Dow Jones Industrial Average (DJIA) index

“ The preprocessing for the text data involved three parts: removing stop words, stemming, and tokenizing. I first removed common words like "the" and "a". This also included removing non-standard english words, such as emojis and words in other languages. The next step was performing stemming, which turns words like "running" to "run". From there, I tried many methods of tokenizing the corpus, including bag of words, term frequency–inverse document frequency, and using bigrams and unigrams. For this particular dataset, a bag of words model using unigrams worked the best.

Using our data from the bag of words model, I merged each day with whether or not the Dow Jones Industrial Average (DJIA) index will move up or down for the next day. I thought the movement of the DJIA would be able to capture the health of the overall economy while not giving too much information into how tech stocks are specifically doing...

The top words associated with DJIA moving up are "gadafi", "Iran" and "sanctions". ”

16.2.2 Time Series in Pytorch | Introductory Example

Q: (cloze) Neural networks can be constructed using the `torch.nn` package.

Q: Import the package for constructing neural networks in PyTorch.

```
import torch.nn as nn
```

Q: (cloze) Seaborn comes with built-in datasets.

Q: Load seaborn's flights dataset.

```
flight_data = sns.load_dataset("flights")
```

Q: Why must time series data be scaled for sequence predictions?

When a network is fit on unscaled data, it is possible for large inputs to slow down the learning and convergence of your network and in some cases prevent the network from effectively learning your problem.

Q: sklearn import for scaling data?

```
from sklearn.preprocessing import MinMaxScaler
```

16.3 Daily Schedule (when free)

Lex Fridman posted a video today (2020 Aug 27), detailing [his daily](#) schedule outside of work commitments. The stages are as follows:

- (1) Morning (2) Mantra (3) Deep work I (4) Medium depth (5) Exercise
- (6) Shower (7) Deep work II (8) Eat (9) Shallow work (10) Paper reading

I'll go into a bit more depth on the sections below and try to implement some of these practices in my daily schedule to see what I can learn from Lex.

1. Morning
2. Mantra / Daily Review
 - Rules & constraints
 - Gratitude
 - Long-term goals
 - Short-term goals
 - Core principles
3. Deep Work I (4 hr)
4. Social Media and Music Practice
5. Exercise

Exercise every day, no matter what. Even if injured, find a body part and exercise it.
6. Shower
7. Deep Work II (4 hr)
8. Eat

2k calories, all your nutrients vegetables, protein, and fat
9. Shallow Work (1 hr)

10. Paper Reading (1 hr)

Chapter 17 2020

10-18-2020